

CS3103: Operating Systems

Spring 2021

Programming Assignment 1

1 Introduction

The purpose of this assignment is to help you better understand process management and system calls in Linux environment. To achieve this goal, you will use C/C++ to implement a **mini shell** that supports running programs not only in the foreground but also in the background: if a program is designated to run in the foreground, then your shell should load and run the program, and wait for the program to terminate before available for the next command line; in contrast, the shell does not wait for the program to terminate before awaiting the next command line if the program runs in the background.

2 Requirements

2.1 ~~Prompt for User Input~~

Your mini shell should display the prompt for user input. For example, if running your mini shell in the Linux terminal,

```
$. / shell
```

it should prompt

```
sh >
```

indicating being ready for user input.

2.2 ~~Foreground Execution of a Program~~

Your mini shell should allow a program to run in the foreground, which has the following form:

```
sh >fg [path of executable file] [a list of arguments]
```

For example, if you want to run the test program mentioned in Section 4 in the foreground, you can type:

```
sh >fg ./test hello 2 5
```

where `./test` is the path of the test program, and **hello 2 5** are the arguments. When running a process in the foreground, the shell is unable to accept the next command from the user until the foreground process terminates. At any time, at most **ONE** process can run in the foreground.

2.3 Background Execution of Programs

The shell should also allow programs to run in the background:

```
sh >bg [path of executable file] [a list of arguments]
```

Unlike the **fg** command, the shell is able to accept the next command immediately after executing the **bg** command. Below is an example:

```
sh >bg ./test hello 2 5
sh >
```

Also note that the shell should allow **ANY** number of processes to exist in the background simultaneously.

~~2.4 Completion of Programs~~

When a foreground process or a background process completes and exits normally, your shell should display a message. For example:

```
sh >fg ./test hello 2 5
hello
hello
hello
hello
hello
Process 4096 completed
sh >
```

~~2.5 Ctrl-C and Ctrl-Z~~

When your mini shell program is running in the Linux terminal, it will receive a SIGINT (SIGTSTP) signal if typing Ctrl-C (Ctrl-Z) on the keyboard. Your shell program is expected to handle the signal and terminate (stop) the foreground process. If there is no foreground process running in the shell, then the signal should have no effect.

For each process that is terminated (stopped) by Ctrl-C (Ctrl-Z), the shell should display the corresponding message. For example, if a process with PID = 4096 is terminated by Ctrl-C, then the shell is expected to display the following message:

```
sh >Process 4096 terminated
```

Similarly, if a process with PID = 4096 is stopped by Ctrl-Z, then the shell is expected to display the following information:

```
sh >Process 4096 stopped
```

2.6 List All Processes

The **list** command should list the information of all processes that are running in the background as well as all the stopped processes. The information should include PID, state, path, and arguments. Note that terminated processes are **NOT** supposed to appear in the list. Therefore, you need to reap the process if it completes and exits normally, or it is terminated by Ctrl-C. Otherwise, the **list** command may behave unexpectedly. Here is an example with correct behavior:

```
sh >bg ./test test1 2 5
sh >bg ./test test2 5 10
sh >list
4096: running ./test test1 2 5
4097: running ./test test2 5 10
sh >fg ./test test3 2 5
Process 4098 terminated
sh >fg ./test test4 5 10
Process 4099 stopped
sh >list
4096: running ./test test1 2 5
4097: running ./test test2 5 10
4099: stopped ./test test4 5 10
sh >
```

~~2.7 Exit the Shell~~

The **exit** command should cause the shell program to exit. But before that, make sure that all processes have been terminated and reaped. Here is an example where there are two processes in the background before executing **exit**:

```
sh >exit
Process 4096 terminated
Process 4097 terminated
$
```

2.8 Handling Unexpected Inputs

You can never make any assumption about what users will input in your shell. For example, a user may input an empty line, an invalid command, or a wrong path of the executable file. You are encouraged to come up with these unexpected inputs as many as possible and handle them properly.

3 Hints

- Study the man pages of the system calls used in your program. For example, the following command displays the man pages of `kill()`:

```
$ man 2 kill
```
- Use `fork()` and `execvp()` so that the parent process accepts user input and the child process executes foreground/background processes.
- If a process is going to run in the foreground, use `sigsuspend()` or `sleep()` in a busy loop until the foreground process terminates.
- Some signals are useful in this assignment. For example, typing `Ctrl-C` will trigger a `SIGINT` signal automatically, and typing `Ctrl-Z` will trigger a `SIGTSTP` signal automatically. Besides, sending a `SIGTERM` signal can terminate a process. Another signal that you might be interested in is the `SIGCHLD` signal, which will be sent to the parent process by the kernel whenever a child process stops or terminates. Your program needs to handle these signals correctly. To handle signals, use `signal()` to bind a signal to the corresponding handler that you implement. By the way, `kill()` is a useful system call for sending a signal (don't get confused by its name).
- When a process terminates for any reason, the process becomes a "zombie" until it is reaped by its parent. To reap a terminated child process (if any), use `waitpid(-1, &status, WNOHANG | WUNTRACED)`. This system call will return immediately, with a return value 0 if none of the children has stopped or terminated, or with a return value equal to the `PID` of one of the stopped or terminated children. When the parent reaps the terminated child, the kernel passes the child's exit status to the parent and then discards the terminated process, at which point it ceases to exist. Also note that the exit status may be useful when handling the `SIGCHLD` signal.
- When you run your shell from the standard Linux shell, your shell is running in the foreground process group of the standard Linux shell. If your shell then creates a child process, by default that child will also be a member of the foreground process group. Since typing `Ctrl-C` sends a `SIGINT` to every process in the foreground group, typing `Ctrl-C` will send a `SIGINT` to your shell, as well as to every process that your shell created. This is obviously not the desired behavior, since we hope only our shell can receive this signal and handle the signal properly. Here is the workaround: After the `fork()`, but before the `execvp()`, the child process should call `setpgid(0, 0)`, which puts the child in a new process group whose group ID is identical to the child's `PID`. This ensures that there will be only one process, your shell, in the foreground process group of the standard Linux shell. When you type `Ctrl-C`, your own shell should catch the resulting `SIGINT` and handle it properly.
- When you use `fork()`, it is important that you do not create a fork bomb, which easily eats up all the resources allocated to you. If this happens, you can try to use the command "kill" to terminate your processes (<http://cslab.cs.cityu.edu.hk/supports/unix-startup-guide>). However, if you cannot log into your account any more, you need to ask CSLab for help to kill your processes.

4 Helper Programs

4.1 args.cpp

This example program shows how to read a line from terminal, as well as parsing (cutting) the string using the `strtok()` function. To compile the program, use the following command:

```
$ g++ args.cpp -lreadline -o args
```

4.2 test.cpp

This program can be used to test your shell. It takes three arguments: the first argument is a single word to be displayed repeatedly; the second argument is the number of seconds between two consecutive displays of the word; the last argument is the number of times the word to be displayed. For example, the following command displays the word “running” 5 times in 2-second interval:

```
$ ./test running 2 5
```

5 Marking

Important Note: Your program will be tested on our CSLab Linux servers (cs3103-01, cs3103-02, cs3103-03). You should describe clearly how to compile and run your program in the text file. **If an executable file cannot be generated and running successfully on our Linux servers, it will be considered as unsuccessful.**

- **fg:** 20%
- **bg:** 20%
- **Ctrl-C:** 10%
- **Ctrl-Z:** 10%
- **list:** 10%
- **exit:** 10%
- **Reaping terminated processes correctly:** 10%
- **Handling unexpected inputs:** 5%
- **Programming style and in-program comments:** 5%

6 Submission

- This assignment is to be done individually or by a group of two students. You are encouraged to discuss the high-level design of your solution with your classmates but you **must** implement the program on your own. Academic dishonesty such as copying another student’s work or allowing another student to copy your work, is regarded as a serious academic offence.
- Each submission consists of two files: a source program file (.cpp file) and a text file containing user guide, if necessary, and all possible outputs produced by your program (.txt file).
- Write down your name(s), eid(s) and student ID(s) in the first few lines of your program as comment.
- Use your student ID(s) to name your submitted files, such as 5xxxxxxx.cpp, 5xxxxxxx.txt for individual submission, or 5xxxxxxx_5yyyyyyy.cpp, 5xxxxxxx_5yyyyyyy.txt for group submission. Only ONE submission is required for each group.
- Submit the files to Canvas.
- The deadline is **11:00am, 18-FEB-2021 (Thursday)**. No late submission will be accepted.

7 Questions?

- This is not a programming course. You are encouraged to debug the program on your own first.
- If you have any questions, please submit your questions to Mr ZHOU Zikang via the Discussion board “Programming Assignment #1”.
- To avoid possible plagiarism, do not post your source code on the Discussion board.
- If necessary, you may also contact Mr ZHOU Zikang at zikanzhou2-c@my.cityu.edu.hk.